

CS100 Lecture 2

Variables I and Arithmetic Types

Contents

- Variable declaration
- Arithmetic types
 - Bits and bytes
 - Integer types
 - Real floating types
 - Character types
 - Boolean type

Variable declaration

Type of a variable

Every variable in C has a type.

- The type is **fully deterministic** and **cannot be changed**.
- The type is **known even when the program is not run**.
 - $\Leftrightarrow$ The type is known at **compile-time**.
 - $\Leftrightarrow$ C is **statically-typed**¹. $\Leftrightarrow$ C has a **static type system**.
 - In contrast, Python is **dynamically-typed**.

Statically-typed vs dynamically-typed

Python: dynamically typed

```
a = 42          # Type of a is int.  
a = "hello"     # Type of a becomes str.
```

C: statically-typed

```
int a = 42;    // Type of a is int.  
a = "hello";  // Error! Types mismatch!
```

The type of a variable

- can be changed, and
- is not necessarily known until we run the program.

The type of a variable

- is explicitly written on declaration, and
- is known at compile-time, and
- cannot be changed.

A type-related error in C is (*usually*) a **compile error**:

- It stops the compiler. The executable will not be generated.

Declare a variable

To declare a variable, we need to specify its **type** and **name**.

```
Type name;
```

Example:

```
int x;    // Declares a variable named `x`, whose type is `int`.  
double y; // Declares a variable named `y`, whose type is `double`.
```

We may declare multiple variables of a same type in one declaration statement, separated by `,`:

```
int x, y; // Declares two variables `x` and `y`, both having type `int`.
```

Declare a variable

A **variable declaration** can be placed

- inside a function, which declares a **local variable**, or
- outside of any functions, which declares a **global variable**.

```
#include <stdio.h>

int x, y; // global variables

int main(void) {
    scanf("%d%d", &x, &y);
    printf("%d\n", x + y);
}
```

```
#include <stdio.h>

int main(void) {
    // local variables in `main`
    int x, y;
    scanf("%d%d", &x, &y);
    printf("%d\n", x + y);
}
```

Local variables vs global variables

Which one do you prefer?

```
#include <stdio.h>

int x, y; // global variables

int main(void) {
    scanf("%d%d", &x, &y);
    printf("%d\n", x + y);
}
```

```
#include <stdio.h>

int main(void) {
    // local variables in `main`
    int x, y;
    scanf("%d%d", &x, &y);
    printf("%d\n", x + y);
}
```

What are these variables used for?

```
#include <stdio.h>
// Other #includes

int x, y; // What are these two variables used for?

int moveSpaceShuttle(SpaceShuttle *shuttle, Coord to, Vehicle *by) {
    // 109 lines
}
int makePreparations(Environment *env, Task tasks[], Time time) {
    // 73 lines
}
LaunchResult launchSpaceShuttle(SpaceShuttle *shuttle, Task tasks[]) {
    // 35 lines
}
// Other 136 functions, 3325 lines in total
int main(void) {
    // 120 lines
}
```

Readability matters

[Best practice] Declare the variable when you first use it!

- If the declaration and use of the variable are too separated, it will become much more difficult to figure out what they are used for as the program goes longer.

[Best practice] Use meaningful names!

- The program would be a mess if polluted with names like `a`, `b`, `c`, `d`, `x`, `y`, `cnt`, `cnt_2`, `flag1`, `flag2`, `flag3` everywhere.
- Use meaningful names: `sumOfScore`, `student_cnt`, `open_success`, ...

Readability is very important. Many students debug day and night simply because their programs are not human-readable.

Use of global variables

One reason for using global variables is to have them shared between functions:

```
int input;
void work(void) {
    printf("%d\n", input);
}
int main(void) {
    scanf("%d", &input);
    work();
}
```

```
void work(void) {
    // Error: `input` was not declared
    // in this scope.
    printf("%d\n", input);
}
int main(void) {
    int input;
    scanf("%d", &input);
    work();
}
```

⇒ More about scopes and name lookup in later lectures / recitations.

Initialize a variable

A variable can be **initialized** on declaration.

```
int x = 42; // Declares the variable `x` of type `int`,  
           // and initializes its value to 42.  
int a = 0, b, c = 42; // Declares three `int` variables, with `a` initialized  
                     // to 0, `c` initialized to 42, and `b` uninitialized.
```

This is syntactically **different** (though seems equivalent) to

```
int x; // Declares `x`, uninitialized.  
x = 42; // Assigns 42 to `x`.
```

[Best practice] Initialize the variable if possible. Prefer initialization to later assignment.

⇒ More on initialization in later lectures.

Arithmetic types

Integer types

Is `int` equivalent to $\mathbb{Z}$?

- Is there a limitation on the numbers that `int` can represent?

Integer types

Is `int` equivalent to $\mathbb{Z}$?

- Is there a limitation on the numbers that `int` can represent?

Experiment:

```
#include <stdio.h>

int main(void) {
    int x = 1;
    while (1) {
        printf("%d\n", x);
        x *= 2; // x = x * 2
        getchar();
    }
}
```

- On 64-bit Ubuntu 22.04 and compiled with GCC 13, after printing `1073741824` (2^{30}), the output becomes negative, and then `0`.

```
1073741824
-2147483648
0
0
```

Bits and bytes

Information is stored in computers in **binary**.

- $42_{\text{ten}} = 101010_{\text{two}}$.

A **bit** is either 0 or 1.

- The binary representation of 42 consists of 6 bits.

A **byte** is 8 bits ² grouped together like 10001001.

- At least 1 byte is needed to store 42.
- At least 3 bytes are needed to store $142857_{\text{ten}} = 100010111000001001_{\text{two}}$

Bits and bytes

A 32-bit number: $2979269462_{\text{ten}} = 10110001100101000000101101010110_{\text{two}}$.



Suppose now we have n bits.

- How many different values can be represented?
- What is the largest integer that can be represented?
- How do we represent negative numbers? Non-integer values? ...

Bits and bytes

Suppose now we have n bits.

- How many different values can be represented?
 - 2^n .
- What is the largest integer that can be represented?
 - $2^n - 1 = \underbrace{111 \dots 1}_n_{\text{two}}$.
- How do we represent negative numbers? Non-integer values? ...
 - There are several different **signed number representations**, among which **two's complement** is widely used.
 - About floating-point numbers: **IEEE754**
 - Details are not covered in CS100.

Integer types

An integer type in C is either **signed** or **unsigned**, and has a **width** denoting the number of bits that can be used to represent values.

Suppose we have an integer type of n bits in width.

- If the type is **signed** ³, the range of values that can be represented is $[-2^{n-1}, 2^{n-1} - 1]$.
- If the type is **unsigned**, the range of values that can be represented is $[0, 2^n - 1]$.

Integer types

(signed)
short (int)

unsigned
short (int)

signed / int /
signed int

unsigned (int)

(signed) long (int)

unsigned long (int)

(signed) long long (int)

unsigned long long (int)

Integer types

- The keyword `int` is optional in types other than `int` :
 - e.g. `short int` and `short` name the same type.
 - e.g. `unsigned int` and `unsigned` name the same type.
- "Unsigned-ness" needs to be written explicitly: `unsigned int`, `unsigned long`, ...
- Types without the keyword `unsigned` are signed by default:
 - e.g. `signed int` and `int` name the same type.
 - e.g. `signed long int`, `signed long`, `long int` and `long` name the same type.

Width of integer types

type	width (at least)	width (usually)
short	16 bits	16 bits
int	16 bits	32 bits
long	32 bits	32 or 64 bits
long long	64 bits	64 bits

- A signed type has the same width as its unsigned counterpart.
- It is also guaranteed that $\text{sizeof}(\text{short}) \leq \text{sizeof}(\text{int}) \leq \text{sizeof}(\text{long}) \leq \text{sizeof}(\text{long long})$.
 - $\text{sizeof}(T)$ is the number of bytes that T holds.

Implementation-defined behaviors

The standard states that the exact width of the integer types is **implementation-defined**.


- **Implementation:** The compiler and the standard library.
- An implementation-defined behavior depends on the compiler and the standard library, and is often also related to the hosted environment (e.g. the operating system).

Which one should I use?

`int` is the most optimal integer type for the platform.

- Use `int` for integer arithmetic by default.
- Use `long long` if the range of `int` is not large enough.
- Use smaller types (`short` , or even `unsigned char`) for memory-saving or other special purposes.
- Use `unsigned` types for special purposes. We will see some in later lectures.

Which one is the real world, the integer types or $\mathbb{Z}$?



Real floating types

"Floating-point": The number's radix point can "float" anywhere to the left, right, or between the significant digits of the number.

Real floating-point types can be used to represent *some* real values.

- Real floating-point types $\neq \mathbb{R}$.

Real floating types

C has three types for representing real floating-point values:

- `float` : single precision. Matches [IEEE754 binary32 format](#) if supported.
- `double` : double precision. Matches [IEEE754 binary64 format](#) if supported.
- `long double` : extended precision. A floating-point type whose precision and range are at least as good as those of `double`.

Details of IEEE754 formats are not required in CS100.

Which one should I use?

Use `double` for real floating-point arithmetic by default.

- In some cases the precision of `float` is not enough.
- Don't worry about efficiency! `double` arithmetic is not necessarily slower than `float`.

Do not use floating-point types for integer arithmetic!

scanf/printf

type	format specifier
short	%hd
int	%d
long	%ld
long long	%lld

type	format specifier
unsigned short	%hu
unsigned	%u
unsigned long	%lu
unsigned long long	%llu

- %f for float, %lf for double, and %Lf for long double.

Exercise

Write the "A+B" program for real numbers. Which type do you decide to use? How do you read and print the values?

Exercise

Write the "A+B" program for real numbers. Which type do you decide to use? How do you read and print the values?

```
#include <stdio.h>

int main(void) {
    double a, b;
    scanf("%lf%lf", &a, &b);
    printf("%lf\n", a + b);
    return 0;
}
```

Character types

The C standard provides three **different** character types: `signed char`, `unsigned char` and `char`.

Let $T \in \{ \text{signed char}, \text{unsigned char}, \text{char} \}$. It is guaranteed that

`1 == sizeof(T) <= sizeof(short) <= sizeof(int) <= sizeof(long) <= sizeof(long long)`.

- **`T` takes exactly 1 byte.**

Question: What is the valid range of `signed char`? `unsigned char`?

Character types

Question: What is the valid range of `signed char` ? `unsigned char` ?

- `signed char` : $[-128, 127]$.
- `unsigned char` : $[0, 255]$.

What? A character is an integer?

ASCII (American Standard Code for Information Interchange)


ASCII assigns a small integer code to each character it includes.

Decimal	Hex	Char	Decimal	Hex	Char	Decimal	Hex	Char	Decimal	Hex	Char
0	0	[NULL]	32	20	[SPACE]	64	40	@	96	60	`
1	1	[START OF HEADING]	33	21	!	65	41	A	97	61	a
2	2	[START OF TEXT]	34	22	"	66	42	B	98	62	b
3	3	[END OF TEXT]	35	23	#	67	43	C	99	63	c
4	4	[END OF TRANSMISSION]	36	24	\$	68	44	D	100	64	d
5	5	[ENQUIRY]	37	25	%	69	45	E	101	65	e
6	6	[ACKNOWLEDGE]	38	26	&	70	46	F	102	66	f
7	7	[BELL]	39	27	'	71	47	G	103	67	g
8	8	[BACKSPACE]	40	28	(	72	48	H	104	68	h
9	9	[HORIZONTAL TAB]	41	29	)	73	49	I	105	69	i
10	A	[LINE FEED]	42	2A	*	74	4A	J	106	6A	j
11	B	[VERTICAL TAB]	43	2B	+	75	4B	K	107	6B	k
12	C	[FORM FEED]	44	2C	,	76	4C	L	108	6C	l
13	D	[CARRIAGE RETURN]	45	2D	-	77	4D	M	109	6D	m
14	E	[SHIFT OUT]	46	2E	.	78	4E	N	110	6E	n
15	F	[SHIFT IN]	47	2F	/	79	4F	O	111	6F	o
16	10	[DATA LINK ESCAPE]	48	30	0	80	50	P	112	70	p
17	11	[DEVICE CONTROL 1]	49	31	1	81	51	Q	113	71	q
18	12	[DEVICE CONTROL 2]	50	32	2	82	52	R	114	72	r
19	13	[DEVICE CONTROL 3]	51	33	3	83	53	S	115	73	s
20	14	[DEVICE CONTROL 4]	52	34	4	84	54	T	116	74	t
21	15	[NEGATIVE ACKNOWLEDGE]	53	35	5	85	55	U	117	75	u
22	16	[SYNCHRONOUS IDLE]	54	36	6	86	56	V	118	76	v
23	17	[ENG OF TRANS. BLOCK]	55	37	7	87	57	W	119	77	w
24	18	[CANCEL]	56	38	8	88	58	X	120	78	x
25	19	[END OF MEDIUM]	57	39	9	89	59	Y	121	79	y
26	1A	[SUBSTITUTE]	58	3A	:	90	5A	Z	122	7A	z
27	1B	[ESCAPE]	59	3B	;	91	5B	[	123	7B	{
28	1C	[FILE SEPARATOR]	60	3C	<	92	5C	\	124	7C	
29	1D	[GROUP SEPARATOR]	61	3D	=	93	5D	]	125	7D	}
30	1E	[RECORD SEPARATOR]	62	3E	>	94	5E	^	126	7E	~
31	1F	[UNIT SEPARATOR]	63	3F	?	95	5F	_	127	7F	[DEL]

ASCII (American Standard Code for Information Interchange)

A character is represented in computers as its [ASCII code](#), which is a small integer.

- We only consider the so-called *ASCII characters* here.



The image shows two side-by-side windows from a development environment. The left window, titled 'C source #1', displays a C program in a text editor with a Vim interface. The code is as follows:

```
1 int main(void) {  
2     char x = 'a';  
3 }
```

The right window, titled 'x86-64 gcc 13.2 (Editor #1)', shows the assembly code generated by the compiler. The assembly is for the 'main' function and includes instructions for stack frame setup, storing the character 'a' (ASCII 97) to memory, and returning. The instruction on line 4 is highlighted with a yellow background, and the value '97' is enclosed in a red box:

```
1 main:  
2     push    rbp  
3     mov     rbp, rsp  
4     mov     BYTE PTR [rbp-1], 97  
5     mov     eax, 0  
6     pop     rbp  
7     ret
```

A character is **nothing but** an integer! In C, there is no "conversion" between characters and ASCII code!

ASCII (American Standard Code for Information Interchange)

Important things to remember:

- ['0' , '9'] = [48, 57].
- ['A' , 'Z'] = [65, 90].
- ['a' , 'z'] = [97, 122].

Example: Given a lowercase letter, return its uppercase form.

```
char to_uppercase(char x) {  
    return x - 32;  
}
```

[Best practice] Avoid magic numbers

What is the meaning of `32` here? $\Rightarrow$ a magic number.

```
char to_uppercase(char x) {  
    return x - 32;  
}
```

Write it in a more human-readable way:

```
char to_uppercase(char x) {  
    return x - ('a' - 'A');  
}
```

Escape sequence

Some special characters are not directly representable: newline, tab, quote, ...

We use [escape sequences](#), e.g.

escape sequence	description
<code>\'</code>	single quote
<code>\"</code>	double quote
<code>\\</code>	backslash

escape sequence	description
<code>\n</code>	newline
<code>\r</code>	carriage return
<code>\t</code>	horizontal tab

Character types

`char`, `signed char` and `unsigned char` are **three different types**.

- Whether `char` is signed or unsigned is **implementation-defined**.
- If `char` is signed (unsigned), it represents the same set of values as the type `signed char` (`unsigned char`), but **they are not the same type**.
 - In contrast, `T` and `signed T` are the same type for $T \in \{ \text{short}, \text{int}, \text{long}, \text{long long} \}$.


Character types

For almost all cases, use `char` (or, sometimes `int`) to represent characters.

`signed char` and `unsigned char` are used for other purposes.

To read/print a `char` using `scanf` / `printf` , use `%c` .

Sad story: Handling non-ASCII characters? ...



```
Developer Command Prompt for VS 2022
*****
** Visual Studio 2022 Developer Command Prompt v17.2.6
** Copyright (c) 2022 Microsoft Corporation
*****

C:\Program Files\Microsoft Visual Studio\2022\Community>cd C:\Users\timur\print_test

C:\Users\timur\print_test>cl.exe main.cpp
Microsoft (R) C/C++ Optimizing Compiler Version 19.32.31332 for x86
Copyright (C) Microsoft Corporation. All rights reserved.

main.cpp
C:\Program Files\Microsoft Visual Studio\2022\Community\VC\Tools\MSVC\14.32.31326\include\ostream(770): warning C4530: C++ exception handler used, but unwind semantics are not enabled. Specify /EHsc
main.cpp(4): note: see reference to function template instantiation 'std::basic_ostream<char,std::char_traits<char>> &std::operator <<<std::char_traits<char>>(std::basic_ostream<char,std::char_traits<char>> &,const char *)' being compiled
Microsoft (R) Incremental Linker Version 14.32.31332.0
Copyright (C) Microsoft Corporation. All rights reserved.

/out:main.exe
main.obj

C:\Users\timur\print_test>main.exe
Привет, мир!

C:\Users\timur\print_test>
```

```
main.cpp - Notepad
File Edit Format View Help

#include <iostream>

int main() {
    std::cout << "Привет, мир!\n";
}

Ln 2, Col 1 100% Unix (LF) UTF-8
```


Sad story: Handling non-ASCII characters? ...

Even though the standard provides `wchar_t`, `char8_t` (since C23), `char16_t` and `char32_t` to handle wide/Unicode characters, there are still a lot of problems.

C++23 has some improvement.

That's why Python people laugh at us ...

**Cppcon** | 2019
The C++ Conference | cppcon.org



Corentin Jabot

Text

Video Sponsorship Provided By:




ALL PROGRAMMING LANGUAGES SUPPORTING UNICODE...

- Swift
- Java/JavaScript
- Python 3
- Safe Rust
- D
- Perl
- PHP
- C#

... EXCEPT

- C
- C++

23



Boolean type: `bool` (since C99)

A type that represents true/false, 0/1, yes/no, ...

To access the name `bool`, `true` and `false`, `<stdbool.h>` is needed. (until C23)

Example: Define a function that accepts a character and returns whether that character is a lowercase letter.

Before C99, using `int`, `0` and `1`:

```
int is_lowercase(char c) {  
    if (c >= 'a' && c <= 'z')  
        return 1;  
    else  
        return 0;  
}
```

Since C99, using `bool`, `false` and `true`:

```
bool is_lowercase(char c) {  
    if (c >= 'a' && c <= 'z')  
        return true;  
    else  
        return false;  
}
```

Boolean type: `bool` (since C99)

Before C99, using `int`, `0` and `1`:

```
int is_lowercase(char c) {  
    if (c >= 'a' && c <= 'z')  
        return 1;  
    else  
        return 0;  
}
```

Since C99, using `bool`, `false` and `true`:

```
bool is_lowercase(char c) {  
    if (c >= 'a' && c <= 'z')  
        return true;  
    else  
        return false;  
}
```

Both return values can be used as follows:

```
char c; scanf("%c", &c);  
if (is_lowercase(c)) {  
    // do something when c is lowercase ...  
}
```

[Best practice] Simplify your code

Just return the result of the condition expression.

```
int is_lowercase(char c) {  
    return c >= 'a' && c <= 'z';  
}
```

```
bool is_lowercase(char c) {  
    return c >= 'a' && c <= 'z';  
}
```

We will introduce the operators (`&&` , `<=` , `>=`) involved here in later lectures.

Summary

- Variable declaration
 - Type + name
 - Multiple variables in one declaration statement
 - Global vs local
 - Initialization

Summary

- Arithmetic types

signed
char

unsigned
char

char

bool

(signed)
short (int)

unsigned
short (int)

char8_t

float

signed / int /
signed int

unsigned (int)

char16_t

double

(signed) long (int)

unsigned long (int)

char32_t

long double

(signed) long long (int)

unsigned long long (int)

wchar_t

Summary

- Arithmetic types
 - Width, signed-ness, valid range
 - Which type to choose
 - Characters: ASCII code, escape sequence
 - Boolean

Exercise

Write a simple calculator that handles input of the form `x op y`, where `x` and `y` are floating-point numbers and `op` $\in \{ \text{'+'}, \text{'-'}, \text{'*'}, \text{'/'} \}$. You may use a group of `if - else` statements like this:

```
if (op == '+') {  
    // ...  
} else if (op == '-') {  
    // ...  
} else if (op == '*') {  
    // ...  
} else if (op == '/') {  
    // ...  
} else {  
    // report an error  
}
```

Exercise

Read one character. If it is a digit ('0' to '9'), convert it to an integer and print its value. Otherwise, print Not a digit.

Define bool is_digit(char c) and call it in main. Store the converted value in an int variable and print it with %d.

```
#include <stdio.h>
#include <stdbool.h>

bool is_digit(char c) {
    // Return whether c is a digit.
}

int main(void) {
    // Read c, call is_digit(c), then convert and print.
}
```

Input: 7 Output: 7

Input: A Output: Not a digit

Notes

- ¹ The type of every expression in C is determined at compile-time except for *variable-length arrays* (since C99).
- ² A byte is 8 bits on most platforms, but we do have exceptions: [36-bit computing](#).
- ³ There are several different signed number representations, but all popular machines and almost all compilers use **two's complement**. Before C23 and C++20, the C/C++ standards allow for all possible representations, so the minimal valid range for a n -bit integer is $[-2^{n-1} + 1, 2^{n-1} - 1]$, which is the range for *one's complement* and *sign-and-magnitude*. Since C23 and C++20, the only representation allowed is two's complement, so the valid range is guaranteed to be $[-2^{n-1}, 2^{n-1} - 1]$. In CS100 we still assume that two's complement is used, even though we are based on C17/C++17.